

Experiment 08

Create and manage NoSQL Databases with Cassandra

Problem Statements:

1. Create keyspace : employee
2. Create : emp_table (
 emp_id int,
 name text,
 city text,
 designation text,
 experience float,
 primary key(emp_id)
);
3. Perform following operations on created table:
 - a. Insert rows
 - b. Update rows
 - c. Update rows with upsert
 - d. Retrieve data from table
 - e. Alter table add columns ((email set<text>, expertise list<text>, prev_jobs map<text, int>)
 - f. Insert new rows
 - g. Delete rows and values
4. create table product(
 id uuid,
 name text,
 price float,
 quan int,
 primary key(id)
);
5. Perform following operations on created table
 - a. Insert rows
 - b. Alter table product add (inv_date timestamp, available boolean);
 - c. Insert new rows

Answers:

1. Create keyspace : employee

```
CREATE KEYSPACE employee
WITH replication = {
    'class': 'SimpleStrategy',
    'replication_factor': 1
};
```

```
cqlsh> CREATE KEYSPACE employee
... WITH replication = {
...     'class': 'SimpleStrategy',
...     'replication_factor': 1
... };
cqlsh> |
```

2. Create collection: inventory

```
USE employee;
```

```
CREATE TABLE emp_table (
    emp_id int PRIMARY KEY,
    name text,
    city text,
    designation text,
    experience float
);
```

```
cqlsh> USE employee;
cqlsh:employee> CREATE TABLE emp_table (
...     emp_id int PRIMARY KEY,
...     name text,
...     city text,
...     designation text,
...     experience float
... );
cqlsh:employee> |
```

3.a. Insert rows

```
INSERT INTO emp_table (emp_id, name, city, designation,
experience)
```

```
VALUES (1, 'John Doe', 'New York', 'Engineer', 5.5);
```

```
INSERT INTO emp_table (emp_id, name, city, designation,
experience)
```

```
VALUES (2, 'Jane Smith', 'Los Angeles', 'Manager', 8.2);
```

```
cqlsh:employee> INSERT INTO emp_table (emp_id, name, city, designation, experience)
... VALUES (1, 'John Doe', 'New York', 'Engineer', 5.5);
cqlsh:employee>
cqlsh:employee> INSERT INTO emp_table (emp_id, name, city, designation, experience)
... VALUES (2, 'Jane Smith', 'Los Angeles', 'Manager', 8.2);
cqlsh:employee> |
```

3.b. Update rows

```
UPDATE emp_table
SET city = 'San Francisco'
WHERE emp_id = 1;
```

```
cqlsh:employee> UPDATE emp_table
... SET city = 'San Francisco'
... WHERE emp_id = 1;
cqlsh:employee>
```

3.c. Update rows with upsert

```
INSERT INTO emp_table (emp_id, name, city, designation,
experience)
```

```
VALUES (2, 'Jane Smith', 'Seattle', 'Designer', 4.0);
```

```
cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | experience | name
-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | 5.5       | John Doe
2 | Los Angeles  | Manager    | 8.2       | Jane Smith

(2 rows)
cqlsh:employee> INSERT INTO emp_table (emp_id, name, city, designation, experience)
... VALUES (2, 'Jane Smith', 'Seattle', 'Designer', 4.0);
cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | experience | name
-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | 5.5       | John Doe
2 | Seattle       | Designer    | 4         | Jane Smith

(2 rows)
cqlsh:employee> |
```

3.d. Retrieve data from table

```
SELECT * FROM emp_table;
```

```
SELECT name, city FROM emp_table WHERE emp_id = 1;
```

```
cqlsh:employee> SELECT * FROM emp_table;
```

emp_id	city	designation	experience	name
1	San Francisco	Engineer	5.5	John Doe
2	Seattle	Designer	4	Jane Smith

```
(2 rows)
cqlsh:employee>
cqlsh:employee> SELECT name, city FROM emp_table WHERE emp_id = 1;
```

name	city
John Doe	San Francisco

```
(1 rows)
cqlsh:employee> |
```

3.e. Alter table add columns ((email set<text>, expertise list<text>, prev_jobs map<text, int>)

```
ALTER TABLE emp_table ADD email set<text>;
```

```
ALTER TABLE emp_table ADD expertise list<text>;
```

```
ALTER TABLE emp_table ADD prev_jobs map<text, int>;
```

```
cqlsh:employee> SELECT * FROM emp_table;
```

emp_id	city	designation	experience	name
1	San Francisco	Engineer	5.5	John Doe
2	Seattle	Designer	4	Jane Smith

```
(2 rows)
cqlsh:employee> ALTER TABLE emp_table ADD email set<text>;
cqlsh:employee>
cqlsh:employee> ALTER TABLE emp_table ADD expertise list<text>;
cqlsh:employee>
cqlsh:employee> ALTER TABLE emp_table ADD prev_jobs map<text, int>;
cqlsh:employee> SELECT * FROM emp_table;
```

emp_id	city	designation	email	experience	expertise	name	prev_jobs
1	San Francisco	Engineer	null	5.5	null	John Doe	null
2	Seattle	Designer	null	4	null	Jane Smith	null

```
(2 rows)
cqlsh:employee> |
```

3. f. Insert new rows

```
INSERT INTO emp_table (emp_id, name, city, designation,
experience, email, expertise, prev_jobs)
```

```
VALUES (
```

```
4, 'Michael Brown', 'Austin', 'Analyst', 3.5,
```

```

        {'michael.b@example.com'},
        ['Java', 'Python'],
        {'Google': 2, 'Facebook': 1}
    );

```

```

cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
2 | Seattle       | Designer   | null           | 4          | null      | Jane Smith    | null

(2 rows)
cqlsh:employee> INSERT INTO emp_table (emp_id, name, city, designation, experience, email, expertise, prev_jobs)
... VALUES (
...     4, 'Michael Brown', 'Austin', 'Analyst', 3.5,
...     {'michael.b@example.com'},
...     ['Java', 'Python'],
...     {'Google': 2, 'Facebook': 1}
... );
cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
2 | Seattle       | Designer   | null           | 4          | null      | Jane Smith    | null
4 | Austin        | Analyst    | {'michael.b@example.com'} | 3.5        | ['Java', 'Python'] | Michael Brown | {'Facebook': 1, 'Google': 2}

(3 rows)
cqlsh:employee> |

```

3. g. Delete rows and values

```
DELETE FROM emp_table WHERE emp_id = 2;
```

```

cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
2 | Seattle       | Designer   | null           | 4          | null      | Jane Smith    | null
4 | Austin        | Analyst    | {'michael.b@example.com'} | 3.5        | ['Java', 'Python'] | Michael Brown | {'Facebook': 1, 'Google': 2}

(3 rows)
cqlsh:employee> DELETE FROM emp_table WHERE emp_id = 2;
cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
4 | Austin        | Analyst    | {'michael.b@example.com'} | 3.5        | ['Java', 'Python'] | Michael Brown | {'Facebook': 1, 'Google': 2}

(2 rows)
cqlsh:employee>

```

```
DELETE email['michael.b@example.com']
FROM emp_table WHERE emp_id = 4;
```

```

cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
4 | Austin        | Analyst    | {'michael.b@example.com'} | 3.5        | ['Java', 'Python'] | Michael Brown | {'Facebook': 1, 'Google': 2}

(2 rows)
cqlsh:employee> DELETE email['michael.b@example.com'] FROM emp_table WHERE emp_id = 4;
cqlsh:employee> SELECT * FROM emp_table;

emp_id | city          | designation | email          | experience | expertise | name          | prev_jobs
-----+-----+-----+-----+-----+-----+-----+-----
1 | San Francisco | Engineer   | null           | 5.5        | null      | John Doe      | null
4 | Austin        | Analyst    | null           | 3.5        | ['Java', 'Python'] | Michael Brown | {'Facebook': 1, 'Google': 2}

(2 rows)
cqlsh:employee> |

```

4. create table

```

product(
    id uuid,
    name text,
    price float,
    quan int,

```

```

cqlsh:employee> CREATE TABLE product (
...     id uuid PRIMARY KEY,
...     name text,
...     price float,
...     quan int
... );
cqlsh:employee> SELECT * FROM product;

id | name | price | quan
---+---+---+---
(0 rows)

```

```
        primary key(id)
    );
```

5. a. Insert rows

```
INSERT INTO product (id, name, price, quan)
VALUES (uuid(), 'Tablet', 399.99, 15);
```

```
INSERT INTO product (id, name, price, quan)
VALUES (uuid(), 'Monitor', 199.49, 20);
```

```
cqlsh:employee> SELECT * FROM product;

 id | name | price | quan
-----+-----+-----+-----
(0 rows)
cqlsh:employee> INSERT INTO product (id, name, price, quan)
... VALUES (uuid(), 'Tablet', 399.99, 15);
cqlsh:employee>
cqlsh:employee> INSERT INTO product (id, name, price, quan)
... VALUES (uuid(), 'Monitor', 199.49, 20);
cqlsh:employee> SELECT * FROM product;

 id | name | price | quan
-----+-----+-----+-----
44ea6ac8-8bb9-4394-a37e-540caf54a66f | Tablet | 399.98999 | 15
872cbfd5-23e2-4e27-a2ae-89c20ef39402 | Monitor | 199.49001 | 20
(2 rows)
```

5. b. Alter table product add (inv_date timestamp, available boolean);

```
ALTER TABLE product ADD inv_date timestamp;
```

```
ALTER TABLE product ADD available boolean;
```

```
cqlsh:employee> SELECT * FROM product;
```

id	name	price	quan
44ea6ac8-8bb9-4394-a37e-540caf54a66f	Tablet	399.98999	15
872cbfd5-23e2-4e27-a2ae-89c20ef39402	Monitor	199.49001	20

(2 rows)

```
cqlsh:employee> ALTER TABLE product ADD inv_date timestamp;
```

```
cqlsh:employee> ALTER TABLE product ADD available boolean;
```

```
cqlsh:employee> SELECT * FROM product;
```

id	available	inv_date	name	price	quan
44ea6ac8-8bb9-4394-a37e-540caf54a66f	null	null	Tablet	399.98999	15
872cbfd5-23e2-4e27-a2ae-89c20ef39402	null	null	Monitor	199.49001	20

(2 rows)

5. c. Insert new rows

```
INSERT INTO product (id, name, price, quan, inv_date,  
available)
```

```
VALUES (uuid(), 'Keyboard', 49.99, 50, '2024-11-25  
10:30:00', true);
```

```
INSERT INTO product (id, name, price, quan, inv_date,  
available)
```

```
VALUES (uuid(), 'Mouse', 29.99, 75, '2024-11-25  
12:45:00', false);
```

```
cqlsh:employee> SELECT * FROM product;
```

id	available	inv_date	name	price	quan
44ea6ac8-8bb9-4394-a37e-540caf54a66f	null	null	Tablet	399.98999	15
872cbfd5-23e2-4e27-a2ae-89c20ef39402	null	null	Monitor	199.49001	20

(2 rows)

```
cqlsh:employee> INSERT INTO product (id, name, price, quan, inv_date, available)
```

```
... VALUES (uuid(), 'Keyboard', 49.99, 50, '2024-11-25 10:30:00', true);
```

```
cqlsh:employee> INSERT INTO product (id, name, price, quan, inv_date, available)
```

```
... VALUES (uuid(), 'Mouse', 29.99, 75, '2024-11-25 12:45:00', false);
```

```
cqlsh:employee> SELECT * FROM product;
```

id	available	inv_date	name	price	quan
180d2c76-7da6-45ec-93b0-fe73c57b59d5	False	2024-11-25 07:15:00+0000	Mouse	29.99	75
44ea6ac8-8bb9-4394-a37e-540caf54a66f	null	null	Tablet	399.98999	15
872cbfd5-23e2-4e27-a2ae-89c20ef39402	null	null	Monitor	199.49001	20
8abcde1c-9423-42aa-a9f1-00d801a4ef6e	True	2024-11-25 05:00:00+0000	Keyboard	49.99	50

(4 rows)